

# Turtlebot Seeker–Target Interception in Unknown Cluttered Environments

## with Monte Carlo Localization and Model Predictive Control

Tim, Gabe, Thomas, Ryan, Kush

## Contents

<b>1 Problem Formulation</b>	<b>2</b>
1.1 State, Control, and Map . . . . .	2
1.2 Objective . . . . .	3
<b>2 System Modeling</b>	<b>3</b>
2.1 Seeker Dynamics (Unicycle Model) . . . . .	3
2.2 Target Dynamics (Constant Velocity Model) . . . . .	4
2.3 Map and LIDAR Measurement Model . . . . .	4
2.4 Target Observation Model . . . . .	4
<b>3 Monte Carlo Localization (MCL)</b>	<b>5</b>
3.1 Bayes Filter . . . . .	5
3.2 Particle Filter Approximation . . . . .	5
3.3 Pose Estimate and Covariance . . . . .	6
<b>4 Target State Estimation via Kalman Filter</b>	<b>6</b>
4.1 Dynamics and Observation . . . . .	6
4.2 Kalman Filter Equations . . . . .	6
<b>5 Model Predictive Control (MPC)</b>	<b>6</b>
5.1 Finite-Horizon Optimal Control Problem . . . . .	7
5.2 Constraints . . . . .	7
5.3 Cost Function . . . . .	7
5.4 Approximate Minimum-Time Behavior . . . . .	8
<b>6 Uncertainty-Aware Constraints</b>	<b>8</b>
6.1 Obstacle Inflation via Pose Uncertainty . . . . .	8
6.2 Interception Region with Uncertainty . . . . .	8
6.3 Speed Limits Based on Uncertainty . . . . .	9
<b>7 Closed-Loop Algorithm</b>	<b>9</b>
<b>8 Offline Probabilistic Trajectory Generation</b>	<b>10</b>
8.1 Stochastic Motion Model . . . . .	10
8.2 Sampling-Based Trajectory Visualization . . . . .	11
8.3 Role of Offline Trajectory in the Real-Time System . . . . .	11

<b>9</b>	<b>Integration of Offline Belief Trajectory with Real-Time MPC</b>	<b>12</b>
<b>10</b>	<b>Algorithm: Offline-to-Online Hybrid Planning</b>	<b>12</b>
<b>11</b>	<b>ROS2 Integration with Probabilistic Planner</b>	<b>12</b>
<b>12</b>	<b>ROS2 Implementation Sketch (Python)</b>	<b>13</b>
12.1	Seeker State Subscriber (using AMCL) . . . . .	13
12.2	Target Kalman Filter Node (Skeleton) . . . . .	14
12.3	MPC Node Skeleton . . . . .	16

# 1 Problem Formulation

We consider two differential-drive robots (Turtlebots) in a cluttered, initially unknown environment:

- A *seeker* robot that attempts to intercept the target.
- A *target* robot that moves through the environment.

The environment is partly unknown and is sensed using 2D LIDAR. The seeker builds a map and estimates its own pose using Monte Carlo Localization (MCL). The target pose is estimated via a separate estimator (e.g. Kalman filter) based on relative measurements. A receding-horizon Model Predictive Controller (MPC) continuously recomputes a trajectory to intercept the target in approximately minimum time while avoiding obstacles and respecting dynamics.

## 1.1 State, Control, and Map

Let the true system state at discrete time  $t$  be

$$X_t = (x_t^{\text{seek}}, x_t^{\text{tgt}}, M_t),$$

where

- $x_t^{\text{seek}}$  is the seeker state (pose and velocity).
- $x_t^{\text{tgt}}$  is the target state (position and velocity).
- $M_t$  is the map / occupancy grid (environment).

The control input applied to the seeker is

$$u_t \in \mathcal{U},$$

typically linear acceleration and angular velocity, or directly wheel commands.

The observations are

$$z_t = (z_t^{\text{LIDAR}}, z_t^{\text{tgt}}),$$

where  $z_t^{\text{LIDAR}}$  is a LIDAR scan and  $z_t^{\text{tgt}}$  is a measurement of the target position (e.g. via vision).

The (stochastic) dynamics and observation models are

$$X_{t+1} \sim p(X_{t+1} | X_t, u_t), \quad (1)$$

$$z_t \sim p(z_t | X_t). \quad (2)$$

The optimal control problem in its most general form is a partially observable stochastic control problem (POMDP). Instead of solving the full POMDP, we:

1. Maintain a belief over seeker pose via MCL.
2. Maintain a belief over target state via (E)KF.
3. Use a point estimate plus uncertainty summaries (covariances) inside a receding-horizon MPC.

## 1.2 Objective

Define an interception condition at some time  $T$ :

$$x_T^{\text{seek}} \in \mathcal{X}_{\text{catch}}(x_T^{\text{tgt}}) := \left\{ x : \|p_T^{\text{seek}} - p_T^{\text{tgt}}\| \leq r_{\text{catch}} \right\},$$

where  $p_T^{\text{seek}} \in \mathbb{R}^2$  is the seeker position and  $p_T^{\text{tgt}}$  is the target position.

The ideal objective is to minimize  $T$  (time to intercept) subject to dynamics, actuator limits, and obstacle avoidance. In practice we approximate this using a fixed-horizon MPC with a cost that strongly rewards rapid reduction of distance to the target.

## 2 System Modeling

### 2.1 Seeker Dynamics (Unicycle Model)

We adopt a unicycle model for the seeker. The state is

$$x_t^{\text{seek}} = \begin{bmatrix} p_{x,t} \\ p_{y,t} \\ \theta_t \\ v_t \end{bmatrix},$$

where  $(p_{x,t}, p_{y,t})$  is position in the map frame,  $\theta_t$  is heading, and  $v_t$  is linear speed.

The control is

$$u_t = \begin{bmatrix} a_t \\ \omega_t \end{bmatrix},$$

where  $a_t$  is linear acceleration and  $\omega_t$  is angular velocity.

Continuous-time dynamics:

$$\dot{p}_x = v \cos \theta, \quad (3)$$

$$\dot{p}_y = v \sin \theta, \quad (4)$$

$$\dot{\theta} = \omega, \quad (5)$$

$$\dot{v} = a. \quad (6)$$

With sampling time  $\Delta t$ , the noise-free discrete-time dynamics are

$$p_{x,t+1} = p_{x,t} + \Delta t v_t \cos \theta_t, \quad (7)$$

$$p_{y,t+1} = p_{y,t} + \Delta t v_t \sin \theta_t, \quad (8)$$

$$\theta_{t+1} = \theta_t + \Delta t \omega_t, \quad (9)$$

$$v_{t+1} = v_t + \Delta t a_t. \quad (10)$$

Including process noise  $w_t$ :

$$x_{t+1}^{\text{seek}} = f(x_t^{\text{seek}}, u_t) + w_t.$$

## 2.2 Target Dynamics (Constant Velocity Model)

For the target, we use a simple nearly-constant-velocity model. The state is

$$x_t^{\text{tgt}} = \begin{bmatrix} p_{x,t}^{\text{tgt}} \\ p_{y,t}^{\text{tgt}} \\ v_{x,t}^{\text{tgt}} \\ v_{y,t}^{\text{tgt}} \end{bmatrix}.$$

The discrete-time dynamics are

$$x_{t+1}^{\text{tgt}} = Ax_t^{\text{tgt}} + w_t^{\text{tgt}},$$

where

$$A = \begin{bmatrix} 1 & 0 & \Delta t & 0 \\ 0 & 1 & 0 & \Delta t \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

The process noise  $w_t^{\text{tgt}}$  models random accelerations.

## 2.3 Map and LIDAR Measurement Model

We represent the map as a 2D occupancy grid:

$$M = \{m_i\}_{i=1}^{N_{\text{cells}}}, \quad m_i \in \{0, 1\},$$

where  $m_i = 1$  indicates an occupied cell.

The LIDAR sensor at time  $t$  produces a set of range measurements

$$z_t^{\text{LIDAR}} = \{r_t^k\}_{k=1}^{N_{\text{beams}}}$$

at fixed angles  $\{\phi^k\}$  in the robot frame. Given a hypothesized pose  $(x, y, \theta)$  and map  $M$ , one can raycast to obtain expected ranges  $\hat{r}^k(x, y, \theta, M)$ .

A common Gaussian measurement model is:

$$p(z_t^{\text{LIDAR}} \mid x_t^{\text{seek}}, M) \propto \exp \left( -\frac{1}{2} \sum_k \frac{(r_t^k - \hat{r}^k(x_t^{\text{seek}}, M))^2}{\sigma_r^2} \right).$$

## 2.4 Target Observation Model

Suppose a vision system provides a noisy measurement of the target position in the map frame:

$$z_t^{\text{tgt}} = \begin{bmatrix} z_{x,t}^{\text{tgt}} \\ z_{y,t}^{\text{tgt}} \end{bmatrix} = Hx_t^{\text{tgt}} + v_t^{\text{tgt}},$$

where

$$H = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}, \quad v_t^{\text{tgt}} \sim \mathcal{N}(0, R).$$

### 3 Monte Carlo Localization (MCL)

#### 3.1 Bayes Filter

The belief over the seeker pose is

$$bel(x_t) = p(x_t^{\text{seek}} \mid z_{0:t}, u_{0:t-1}, M).$$

The Bayes filter recursion consists of:

**Prediction:**

$$\overline{bel}(x_t) = \int p(x_t \mid x_{t-1}, u_{t-1}) bel(x_{t-1}) dx_{t-1}.$$

**Update:**

$$bel(x_t) = \eta p(z_t^{\text{LIDAR}} \mid x_t, M) \overline{bel}(x_t),$$

where  $\eta$  is a normalizing constant.

#### 3.2 Particle Filter Approximation

We approximate the belief by a weighted set of particles:

$$bel(x_t) \approx \sum_{i=1}^{N_p} w_t^{(i)} \delta(x_t - x_t^{(i)}),$$

where  $x_t^{(i)}$  is the  $i$ -th particle and  $w_t^{(i)}$  its weight.

#### MCL Algorithm

---

**Algorithm 1** Monte Carlo Localization (one time step)

---

**Require:** particles  $\{x_{t-1}^{(i)}, w_{t-1}^{(i)}\}_{i=1}^{N_p}$ , control  $u_{t-1}$ , LIDAR scan  $z_t^{\text{LIDAR}}$ , map  $M$

**Ensure:** new particles  $\{x_t^{(i)}, w_t^{(i)}\}_{i=1}^{N_p}$

1: **for**  $i = 1$  to  $N_p$  **do**

2:     Sample motion:

$$x_t^{(i)} \sim p(x_t \mid x_{t-1}^{(i)}, u_{t-1})$$

using the unicycle motion model plus noise.

3:     Compute importance weight:

$$w_t^{(i)} \leftarrow p(z_t^{\text{LIDAR}} \mid x_t^{(i)}, M)$$

4: **end for**

5: Normalize weights:

$$w_t^{(i)} \leftarrow \frac{w_t^{(i)}}{\sum_j w_t^{(j)}}$$

6: Resample  $N_p$  particles from  $\{x_t^{(i)}\}$  according to  $\{w_t^{(i)}\}$ .

---

### 3.3 Pose Estimate and Covariance

From the particle set, approximate the mean pose as:

$$\hat{x}_t^{\text{seek}} = \sum_{i=1}^{N_p} w_t^{(i)} x_t^{(i)}.$$

The covariance is

$$P_t^{\text{seek}} = \sum_{i=1}^{N_p} w_t^{(i)} (x_t^{(i)} - \hat{x}_t) (x_t^{(i)} - \hat{x}_t)^{\top}.$$

In practice one often only uses the  $2 \times 2$  position covariance.

## 4 Target State Estimation via Kalman Filter

Because the target model is linear with Gaussian noise, we use a Kalman Filter.

### 4.1 Dynamics and Observation

Dynamics:

$$x_{t+1}^{\text{tgt}} = Ax_t^{\text{tgt}} + w_t^{\text{tgt}}, \quad w_t^{\text{tgt}} \sim \mathcal{N}(0, Q).$$

Observation:

$$z_t^{\text{tgt}} = Hx_t^{\text{tgt}} + v_t^{\text{tgt}}, \quad v_t^{\text{tgt}} \sim \mathcal{N}(0, R).$$

### 4.2 Kalman Filter Equations

Let  $\hat{x}_{t|t}$  be the filtered estimate and  $P_{t|t}$  the corresponding covariance.

#### Prediction

$$\hat{x}_{t|t-1} = A\hat{x}_{t-1|t-1}, \tag{11}$$

$$P_{t|t-1} = AP_{t-1|t-1}A^{\top} + Q. \tag{12}$$

#### Update

$$S_t = HP_{t|t-1}H^{\top} + R, \tag{13}$$

$$K_t = P_{t|t-1}H^{\top}S_t^{-1}, \tag{14}$$

$$\hat{x}_{t|t} = \hat{x}_{t|t-1} + K_t(z_t^{\text{tgt}} - H\hat{x}_{t|t-1}), \tag{15}$$

$$P_{t|t} = (I - K_tH)P_{t|t-1}. \tag{16}$$

The KF provides  $\hat{x}_t^{\text{tgt}}$  and  $P_t^{\text{tgt}}$ .

## 5 Model Predictive Control (MPC)

MPC repeatedly solves a finite-horizon optimal control problem using current state estimates and predicted target motion.

## 5.1 Finite-Horizon Optimal Control Problem

At time  $t$ , let  $\hat{x}_t^{\text{seek}}$  be the seeker pose estimate and  $\hat{x}_t^{\text{tgt}}$  the target state estimate. We define a prediction horizon of length  $N$ . For prediction step  $k = 0, \dots, N$  we have:

- seeker state:  $x_k$ ,
- seeker control:  $u_k$ ,
- predicted target state:  $x_k^{\text{tgt}}$ .

The predicted target states are obtained by propagating the KF estimate:

$$x_{k+1}^{\text{tgt}} = Ax_k^{\text{tgt}}.$$

The seeker dynamics are:

$$x_{k+1} = f(x_k, u_k), \quad k = 0, \dots, N-1,$$

with  $x_0 = \hat{x}_t^{\text{seek}}$ .

## 5.2 Constraints

We impose:

- **Actuator limits:**

$$v_{\min} \leq v_k \leq v_{\max}, \quad a_{\min} \leq a_k \leq a_{\max}, \quad |\omega_k| \leq \omega_{\max}.$$

- **Obstacle avoidance:** for a set of obstacles approximated as circles with centers  $c^o$  and radii  $r^o$ ,

$$\|p_k - c^o\|_2 \geq r_{\text{eff}}^o,$$

where  $r_{\text{eff}}^o$  is an inflated radius (see Section 6).

## 5.3 Cost Function

Define the seeker and target positions:

$$p_k = \begin{bmatrix} p_{x,k} \\ p_{y,k} \end{bmatrix}, \quad p_k^{\text{tgt}} = \begin{bmatrix} p_{x,k}^{\text{tgt}} \\ p_{y,k}^{\text{tgt}} \end{bmatrix}.$$

A typical quadratic stage cost is:

$$\begin{aligned} \ell(x_k, u_k, x_k^{\text{tgt}}) &= (p_k - p_k^{\text{tgt}})^T Q (p_k - p_k^{\text{tgt}}) + q_\theta (\theta_k - \theta_k^{\text{tgt}})^2 \\ &\quad + u_k^T R u_k + (u_k - u_{k-1})^T R_\Delta (u_k - u_{k-1}), \end{aligned} \tag{17}$$

with  $u_{-1}$  taken as the previously applied control. The terminal cost is:

$$\ell_f(x_N, x_N^{\text{tgt}}) = (p_N - p_N^{\text{tgt}})^T Q_f (p_N - p_N^{\text{tgt}}).$$

The total horizon cost is

$$J = \sum_{k=0}^{N-1} \ell(x_k, u_k, x_k^{\text{tgt}}) + \ell_f(x_N, x_N^{\text{tgt}}).$$

## 5.4 Approximate Minimum-Time Behavior

Strict time-optimal control would minimize the interception time  $T$  explicitly. Instead, we approximate minimum time by the choice of cost and constraints.

Let

$$d_k = \|p_k - p_k^{\text{tgt}}\|$$

be the distance between seeker and target. One can add a “progress” term:

$$\ell_{\text{progress}} = -\alpha(d_{k-1} - d_k),$$

which rewards large reductions in distance per step.

An approximate time-to-go can be defined as

$$T_{\text{guess}} = \frac{\|\hat{p}_t^{\text{seek}} - \hat{p}_t^{\text{tgt}}\|}{v_{\max}}.$$

The weighting matrices  $Q$ ,  $Q_f$ , and  $\alpha$  can be adapted as functions of  $T_{\text{guess}}$  to make the controller more aggressive when close to interception.

## 6 Uncertainty-Aware Constraints

The MCL and KF provide covariances  $P_t^{\text{seek}}$  and  $P_t^{\text{tgt}}$  that quantify uncertainty. We use them to adapt constraints.

### 6.1 Obstacle Inflation via Pose Uncertainty

Let  $P_t^{\text{seek}}$  be the  $2 \times 2$  position covariance of the seeker. Let  $\lambda_{\max}(P_t^{\text{seek}})$  be its largest eigenvalue. Define a scalar uncertainty measure:

$$\sigma_t^{\text{seek}} = \sqrt{\lambda_{\max}(P_t^{\text{seek}})}.$$

For an obstacle with base safe radius  $r^o$ , define the inflated radius:

$$r_{\text{eff}}^o = r^o + k_\sigma \sigma_t^{\text{seek}},$$

where  $k_\sigma > 0$  is a tuning parameter.

The deterministic collision avoidance constraint at prediction step  $k$  becomes:

$$\|p_k - c^o\|_2 \geq r_{\text{eff}}^o.$$

As localization improves (smaller  $\sigma_t^{\text{seek}}$ ),  $r_{\text{eff}}^o$  shrinks and the robot can drive closer to obstacles.

### 6.2 Interception Region with Uncertainty

Let the seeker and target positions at step  $k$  be independent Gaussians:

$$p_k^{\text{seek}} \sim \mathcal{N}(\mu_k^{\text{seek}}, \Sigma_k^{\text{seek}}), \quad (18)$$

$$p_k^{\text{tgt}} \sim \mathcal{N}(\mu_k^{\text{tgt}}, \Sigma_k^{\text{tgt}}). \quad (19)$$

The difference

$$d_k = p_k^{\text{seek}} - p_k^{\text{tgt}}$$

is Gaussian with

$$d_k \sim \mathcal{N}(\mu_k^d, \Sigma_k^d),$$

where

$$\mu_k^d = \mu_k^{\text{seek}} - \mu_k^{\text{tgt}}, \quad (20)$$

$$\Sigma_k^d = \Sigma_k^{\text{seek}} + \Sigma_k^{\text{tgt}}. \quad (21)$$

A chance constraint for interception could be

$$\mathbb{P}(\|d_k\| \leq r_{\text{catch}}) \geq 1 - \delta.$$

An approximate deterministic surrogate is

$$\|\mu_k^d\| \leq r_{\text{catch}} - \beta_\delta \sqrt{\lambda_{\max}(\Sigma_k^d)},$$

where  $\beta_\delta$  is chosen from Gaussian tail bounds.

In practice, we may simply set

$$r_{\text{catch,eff}} = r_{\text{catch,base}} + k_{\sigma, \text{catch}} \sqrt{\lambda_{\max}(\Sigma_k^d)}$$

and require  $\|\mu_k^d\| \leq r_{\text{catch,eff}}$ .

### 6.3 Speed Limits Based on Uncertainty

One can also modulate speed based on seeker uncertainty:

$$s_t = \sqrt{\lambda_{\max}(P_t^{\text{seek}})},$$

$$v_{\max}(t) = v_{\max}^{\text{base}} e^{-\alpha s_t},$$

for some  $\alpha > 0$ . The speed constraint becomes

$$|v_k| \leq v_{\max}(t).$$

As MCL converges and uncertainty shrinks,  $v_{\max}(t)$  approaches  $v_{\max}^{\text{base}}$  and the robot can move faster.

## 7 Closed-Loop Algorithm

We now summarize the overall closed-loop algorithm.

---

**Algorithm 2** Closed-Loop MCL + KF + MPC

---

- 1: Initialize MCL particles for seeker pose
- 2: Initialize KF for target state
- 3: **while** experiment running **do**
- 4:   Receive control  $u_{t-1}$  that was sent at previous step
- 5:   Receive new LIDAR scan  $z_t^{\text{LIDAR}}$  and target measurement  $z_t^{\text{tgt}}$
- 6:   **MCL update:**
- 7:      $\text{bel}(x_t^{\text{seek}}) \leftarrow \text{MCL}(\text{bel}(x_{t-1}^{\text{seek}}), u_{t-1}, z_t^{\text{LIDAR}}, M_t)$
- 8:     Extract  $\hat{x}_t^{\text{seek}}, P_t^{\text{seek}}$
- 9:   **KF update:**
- 10:     $(\hat{x}_t^{\text{tgt}}, P_t^{\text{tgt}}) \leftarrow \text{KF}(\hat{x}_{t-1}^{\text{tgt}}, P_{t-1}^{\text{tgt}}, z_t^{\text{tgt}})$
- 11:   **Prediction for horizon:**
- 12:     Predict target states  $x_k^{\text{tgt}}$  for  $k = 0, \dots, N$
- 13:     Set initial seeker state  $x_0 = \hat{x}_t^{\text{seek}}$
- 14:     Compute inflated obstacle radii and interception radius based on  $P_t^{\text{seek}}, P_t^{\text{tgt}}$
- 15:   **MPC:**
- 16:     Solve finite-horizon OCP for  $x_k, u_k$  to minimize cost subject to dynamics and constraints
- 17:     Extract  $u_t = u_0^*$
- 18:   **Apply control:**
- 19:     Send  $u_t$  to seeker robot
- 20:      $t \leftarrow t + 1$
- 21: **end while**

---

## 8 Offline Probabilistic Trajectory Generation

To initialize the real-time interception system, we compute an *offline probabilistic trajectory* from the current seeker pose to a desired goal location. This trajectory is generated using a stochastic motion model together with noise propagation, producing a sequence of waypoints

$$\{(\bar{x}_k, \Sigma_k)\}_{k=0}^{N_{\text{off}}},$$

where  $\bar{x}_k$  is the nominal waypoint pose and  $\Sigma_k$  is the predicted covariance of pose uncertainty at that point. These covariance matrices define confidence ellipses along the trajectory, giving a spatial representation of uncertainty (Figure 1).

### 8.1 Stochastic Motion Model

We assume the seeker obeys the unicycle dynamics

$$x_{k+1} = f(x_k, u_k) + w_k,$$

with process noise

$$w_k \sim \mathcal{N}(0, Q_k).$$

Given a nominal control sequence  $\{u_k\}$  produced by a deterministic planner (e.g. straight-line interpolation, A\* path, RRT), we propagate uncertainty using the first-order linearization

$$A_k = \frac{\partial f}{\partial x}\Big|_{\bar{x}_k, u_k}, \quad B_k = \frac{\partial f}{\partial u}\Big|_{\bar{x}_k, u_k}.$$

The covariance recursion becomes

$$\Sigma_{k+1} = A_k \Sigma_k A_k^\top + Q_k.$$

The result is a *belief trajectory*:

$$\bar{x}_0, \bar{x}_1, \dots, \bar{x}_{N_{\text{off}}}, \quad \Sigma_0, \Sigma_1, \dots, \Sigma_{N_{\text{off}}}.$$

## 8.2 Sampling-Based Trajectory Visualization

In addition to covariance propagation, Monte Carlo realizations

$$x_k^{(i)} = f(x_{k-1}^{(i)}, u_{k-1}) + w_{k-1}^{(i)}, \quad i = 1, \dots, N_s,$$

are generated to visualize dispersion around the nominal path. Plotting these noisy sampled trajectories produces the red ensemble shown in Figure 1. Confidence ellipses, computed from  $\Sigma_k$ , illustrate the predicted uncertainty at each waypoint.

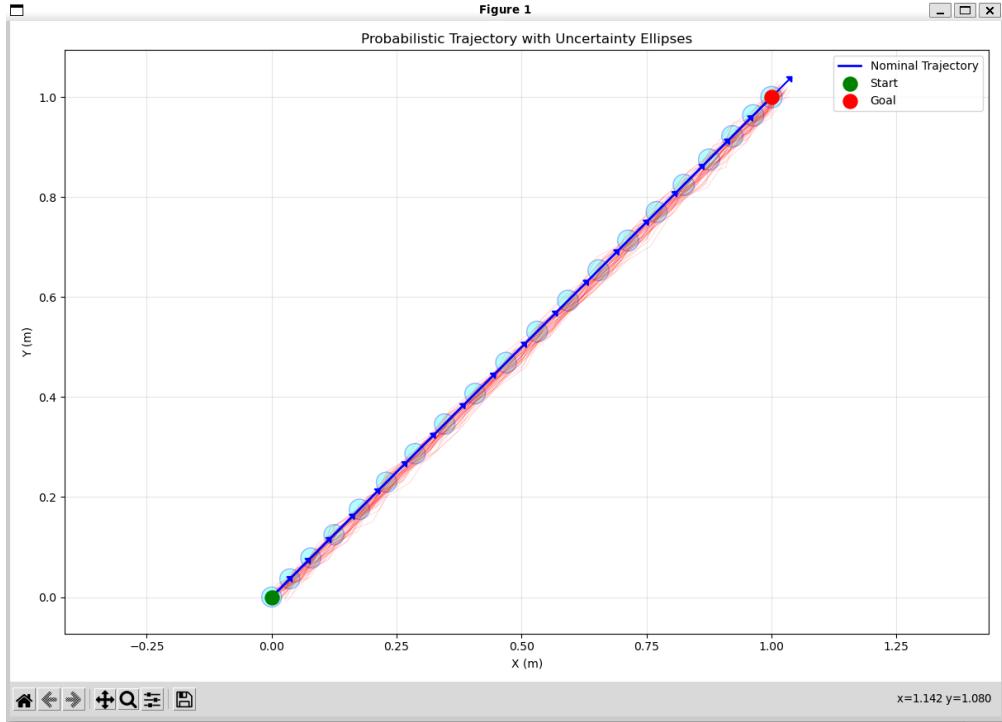


Figure 1: Offline probabilistic trajectory with covariance ellipses and sampled paths.

## 8.3 Role of Offline Trajectory in the Real-Time System

The offline trajectory serves three purposes:

1. It provides an initial **belief-state path** for the MPC to track.
2. It provides **uncertainty priors**  $\Sigma_k$  used to shape:
  - obstacle inflation radii,

- speed limits,
  - catching-radius uncertainty bounds during interception.
3. It initializes the controller before MCL converges, avoiding erratic early-stage control.

During real-time operation, these offline waypoints are *not* followed blindly. Instead, they act as a *warm start* for the adaptive MPC and belief updates, which then continuously re-plan based on real measurements.

## 9 Integration of Offline Belief Trajectory with Real-Time MPC

Once the experiment begins, the belief about the seeker pose is maintained via MCL:

$$b_t(x) = p(x_t^{\text{seek}} \mid z_{0:t}, u_{0:t-1}).$$

Simultaneously, the target is estimated using a Kalman filter.

Let  $\hat{x}_t^{\text{seek}}$  and  $P_t^{\text{seek}}$  denote the filtered mean and covariance of the seeker at time  $t$ . The offline trajectory provides:

$$(\bar{x}_k, \Sigma_k) \quad \text{for } k = 0, \dots, N_{\text{off}}.$$

At each MPC cycle:

1. The controller compares the online belief  $(\hat{x}_t, P_t)$  to the offline predicted belief  $(\bar{x}_{k(t)}, \Sigma_{k(t)})$ .
2. The MPC cost includes a belief-tracking term:

$$\ell_{\text{belief}} = (\hat{x}_t - \bar{x}_{k(t)})^\top Q_b (\hat{x}_t - \bar{x}_{k(t)}).$$

3. The uncertainty  $\Sigma_k$  shapes the online MPC constraints:

$$r_{\text{inflated}} = r_0 + k_\sigma \sqrt{\lambda_{\max}(\Sigma_k)},$$

where  $r_{\text{inflated}}$  is the current safe obstacle buffer.

4. If MCL yields lower covariance ( $P_t \ll \Sigma_k$ ), the MPC automatically becomes more aggressive.
5. If MCL uncertainty grows, the MPC reverts to safer velocities and larger buffers.

Thus, the offline trajectory acts as a prior over feasible robot motion, while the online MPC adaptively re-optimizes the control inputs in real time based on updated belief estimates.

## 10 Algorithm: Offline-to-Online Hybrid Planning

## 11 ROS2 Integration with Probabilistic Planner

The provided ROS2 node integrates the offline trajectory through:

- calls to `plan_probabilistic_trajectory()` which returns waypoints with their covariances,
- publishing noisy (simulated) pose estimates,
- a waypoint-level controller that moves the robot along the nominal offline path while respecting predicted uncertainty.

---

**Algorithm 3** Hybrid Offline Probabilistic Planning + Online MPC

---

- 1: **Offline Phase:**
- 2: Generate nominal controls  $\{u_k\}_{k=0}^{N_{\text{off}}}$
- 3: Propagate states  $\bar{x}_{k+1} = f(\bar{x}_k, u_k)$
- 4: Propagate covariance  $\Sigma_{k+1} = A_k \Sigma_k A_k^\top + Q_k$
- 5: Store  $(\bar{x}_k, \Sigma_k)$
- 6: **Online Phase (Real-Time Loop):**
- 7: **while** robot active **do**
- 8:   MCL update  $\Rightarrow (\hat{x}_t^{\text{seek}}, P_t^{\text{seek}})$
- 9:   KF update  $\Rightarrow (\hat{x}_t^{\text{tgt}}, P_t^{\text{tgt}})$
- 10:   Determine nearest offline waypoint  $(\bar{x}_{k(t)}, \Sigma_{k(t)})$
- 11:   Inflate obstacles using  $\Sigma_{k(t)}$  or  $P_t^{\text{seek}}$
- 12:   Formulate MPC with:
  - belief-tracking cost,
  - target interception cost,
  - uncertainty-dependent constraints.
- 13:   Solve MPC  $\Rightarrow u_t^*$
- 14:   Apply  $u_t^*$  to seeker
- 15: **end while**

---

In the full system, this node is replaced or augmented by:

1. an MCL node producing online beliefs,
2. a target-tracker node producing target beliefs,
3. an MPC node that re-plans dynamically,
4. but retains the offline waypoint covariance sequence as an MPC warm start.

## 12 ROS2 Implementation Sketch (Python)

This section provides minimal ROS2 node skeletons in Python using `rclpy`. They are intended as starting points only.

### 12.1 Seeker State Subscriber (using AMCL)

We assume the standard `amcl` node publishes `/amcl_pose` of type `geometry_msgs/msg/PoseWithCovarianceStamped`.

Listing 1: `seeker_state_node.py`

```
#!/usr/bin/env python3
import rclpy
from rclpy.node import Node
from geometry_msgs.msg import PoseWithCovarianceStamped
from geometry_msgs.msg import Twist

import numpy as np
```

```

class SeekerStateNode(Node):
    def __init__(self):
        super().__init__('seeker_state_node')
        self.pose_sub = self.create_subscription(
            PoseWithCovarianceStamped,
            '/amcl_pose',
            self.pose_callback,
            10
        )
        self.cmd_pub = self.create_publisher(Twist, '/cmd_vel', 10)
        self.latest_pose = None
        self.latest_cov = None

    def pose_callback(self, msg: PoseWithCovarianceStamped):
        self.latest_pose = msg.pose.pose
        self.latest_cov = np.array(msg.pose.covariance).reshape((6, 6))
        # For debugging:
        # self.get_logger().info(f"Pose: {self.latest_pose.position.x}, {self.
        # latest_pose.position.y}")

    def main(args=None):
        rclpy.init(args=args)
        node = SeekerStateNode()
        rclpy.spin(node)
        node.destroy_node()
        rclpy.shutdown()

if __name__ == '__main__':
    main()

```

## 12.2 Target Kalman Filter Node (Skeleton)

Listing 2: target\_kf\_node.py

```

#!/usr/bin/env python3
import rclpy
from rclpy.node import Node
from geometry_msgs.msg import PoseWithCovarianceStamped
from geometry_msgs.msg import PoseStamped
import numpy as np

class TargetKFNodde(Node):
    def __init__(self):
        super().__init__('target_kf_node')

        # state: [px, py, vx, vy]^T
        self.x = np.zeros((4, 1))
        self.P = np.eye(4) * 1.0

        self.dt = 0.1

        self.A = np.array([

```

```

        [1.0, 0.0, self.dt, 0.0],
        [0.0, 1.0, 0.0, self.dt],
        [0.0, 0.0, 1.0, 0.0],
        [0.0, 0.0, 0.0, 1.0],
    ])

self.Q = np.eye(4) * 0.01

self.H = np.array([
    [1.0, 0.0, 0.0, 0.0],
    [0.0, 1.0, 0.0, 0.0],
])
self.R = np.eye(2) * 0.05

self.meas_sub = self.create_subscription(
    PoseStamped,
    '/target_pose_measurement',
    self.measurement_callback,
    10
)

self.est_pub = self.create_publisher(
    PoseWithCovarianceStamped,
    '/target_estimate',
    10
)

self.timer = self.create_timer(self.dt, self.timer_callback)

def measurement_callback(self, msg: PoseStamped):
    z = np.array([[msg.pose.position.x],
                  [msg.pose.position.y]])
    # Update step
    x_pred = self.A @ self.x
    P_pred = self.A @ self.P @ self.A.T + self.Q

    S = self.H @ P_pred @ self.H.T + self.R
    K = P_pred @ self.H.T @ np.linalg.inv(S)

    self.x = x_pred + K @ (z - self.H @ x_pred)
    self.P = (np.eye(4) - K @ self.H) @ P_pred

def timer_callback(self):
    # Prediction-only if no new measurement
    self.x = self.A @ self.x
    self.P = self.A @ self.P @ self.A.T + self.Q

    msg = PoseWithCovarianceStamped()
    msg.header.stamp = self.get_clock().now().to_msg()
    msg.header.frame_id = 'map'
    msg.pose.pose.position.x = float(self.x[0])
    msg.pose.pose.position.y = float(self.x[1])
    # Orientation unused here

```

```

cov = np.zeros((6, 6))
cov[0, 0] = self.P[0, 0]
cov[1, 1] = self.P[1, 1]
msg.pose.covariance = cov.flatten().tolist()
self.est_pub.publish(msg)

def main(args=None):
    rclpy.init(args=args)
    node = TargetKFNode()
    rclpy.spin(node)
    node.destroy_node()
    rclpy.shutdown()

if __name__ == '__main__':
    main()

```

### 12.3 MPC Node Skeleton

This node subscribes to seeker and target estimates and publishes a velocity command. The actual optimization (e.g. using cvxpy) is represented as a placeholder function.

Listing 3: mpc\_node.py (skeleton)

```

#!/usr/bin/env python3
import rclpy
from rclpy.node import Node
from geometry_msgs.msg import PoseWithCovarianceStamped, Twist
import numpy as np

class MPCNode(Node):
    def __init__(self):
        super().__init__('mpc_node')

        self.seeker_sub = self.create_subscription(
            PoseWithCovarianceStamped,
            '/amcl_pose',
            self.seeker_callback,
            10
        )

        self.target_sub = self.create_subscription(
            PoseWithCovarianceStamped,
            '/target_estimate',
            self.target_callback,
            10
        )

        self.cmd_pub = self.create_publisher(Twist, '/cmd_vel', 10)

        self.seeker_pose = None
        self.seeker_cov = None
        self.target_pose = None
        self.target_cov = None

```

```

    self.dt = 0.1
    self.N = 15 # MPC horizon

    self.timer = self.create_timer(self.dt, self.timer_callback)

    def seeker_callback(self, msg: PoseWithCovarianceStamped):
        self.seeker_pose = msg.pose.pose
        self.seeker_cov = np.array(msg.pose.covariance).reshape((6, 6))

    def target_callback(self, msg: PoseWithCovarianceStamped):
        self.target_pose = msg.pose.pose
        self.target_cov = np.array(msg.pose.covariance).reshape((6, 6))

    def timer_callback(self):
        if self.seeker_pose is None or self.target_pose is None:
            return

        # Build initial seeker state x0
        x0 = np.zeros((4, 1))
        x0[0, 0] = self.seeker_pose.position.x
        x0[1, 0] = self.seeker_pose.position.y
        # TODO: extract yaw from quaternion
        # x0[2, 0] = theta
        # x0[3, 0] = v

        # Build target prediction sequence (constant position for now)
        target_seq = []
        px_tgt = self.target_pose.position.x
        py_tgt = self.target_pose.position.y
        for k in range(self.N + 1):
            target_seq.append(np.array([px_tgt, py_tgt]))

        # TODO: compute inflated obstacle radii from seeker_cov, target_cov
        # obstacles = [...]

        # Solve MPC optimization (placeholder)
        v_cmd, w_cmd = self.solve_mpc(x0, target_seq)

        twist = Twist()
        twist.linear.x = float(v_cmd)
        twist.angular.z = float(w_cmd)
        self.cmd_pub.publish(twist)

    def solve_mpc(self, x0, target_seq):
        # Placeholder MPC solver:
        # For now, simply drive straight towards the target.
        px = x0[0, 0]
        py = x0[1, 0]
        tgt = target_seq[0]
        dx = tgt[0] - px
        dy = tgt[1] - py
        dist = np.hypot(dx, dy)
        v_cmd = min(0.5, dist) # saturate speed
        w_cmd = 0.0

```

```
    return v_cmd, w_cmd

def main(args=None):
    rclpy.init(args=args)
    node = MPCNode()
    rclpy.spin(node)
    node.destroy_node()
    rclpy.shutdown()

if __name__ == '__main__':
    main()
```